

Revisiting pitch framing with Bayesian Additive Regression Trees

Sameer K. Deshpande


University of Wisconsin–Madison

EURO 2022

6 July 2022

Pitch framing

- Umpires' ball/strike decisions not deterministic
- Some catchers can influence $\mathbb{P}(\text{strike})$
 - ▶ Good framers “make balls look like strikes”
 - ▶ Good framers “steal strikes”
- Seemingly outsized value relative to # of frameable pitches
 - ▶ Top framers save 15–25 runs on avg. per season
 - ▶ Estimated \$14M/yr in contract value
 - ▶ Only a few frameable pitches per game
- Motivating questions:
 - ▶ How certain are we about the value of framing?
 - ▶ Is estimated framing value stable season-to-season?



Estimating the impact of pitch framing

- Catcher expected to save $\rho \times (p - p_0)$ runs over baseline
 - ▶ ρ : expected value of a called strike relative to ball
 - ▶ p & p_0 : called strike prob. & league-average baseline prob.
 - 😊 Avoids awarding credit for obvious calls & accounts for context
 - 😞 We don't know the exact values ρ , p or p_0 for each pitch

Estimating the impact of pitch framing

- Catcher expected to save $\rho \times (p - p_0)$ runs over baseline
 - ▶ ρ : expected value of a called strike relative to ball
 - ▶ p & p_0 : called strike prob. & league-average baseline prob.
 - 😊 Avoids awarding credit for obvious calls & accounts for context
 - 😞 We don't know the exact values ρ , p or p_0 for each pitch
- Value of a called strike estimated using historic data
 - ▶ ρ : diff. in avg. # runs in scored in half-inning after called ball & strike
 - ▶ Assumption: ρ depends only on count & outs
 - ▶ E.g. value of a strike on 1-1 pitch w/ 1 out is 0.14

Estimating the impact of pitch framing

- Catcher expected to save $\rho \times (p - p_0)$ runs over baseline
 - ▶ ρ : expected value of a called strike relative to ball
 - ▶ p & p_0 : called strike prob. & league-average baseline prob.
 - 😊 Avoids awarding credit for obvious calls & accounts for context
 - 😞 We don't know the exact values ρ , p or p_0 for each pitch
- Value of a called strike estimated using historic data
 - ▶ ρ : diff. in avg. # runs in scored in half-inning after called ball & strike
 - ▶ Assumption: ρ depends only on count & outs
 - ▶ E.g. value of a strike on 1-1 pitch w/ 1 out is 0.14
- Estimated called strike probabilities
 - ▶ $\mathbb{P}(\text{called strike}) = F(\text{location}, \text{personnel}, \text{context})$
 - ▶ context: count, outs, baserunner configuration, ...
 - ▶ Baseline p_0 : average of “counterfactual” catcher probs.
 - ▶ Need to estimate F and get reliable uncertainty estimates

Remaining plan

- D & Wyner (2017)'s hierarchical model for \mathbb{P} (called `strike`)
 - 😊 “Shares statistical strength” between umpires
 - 😞 Heavily over-parametrized & indirect parametrization of `location`
 - 😞 Limited to two-way interactions (e.g. `umpire` \times `count`)

Remaining plan

- D & Wyner (2017)'s hierarchical model for \mathbb{P} (called `strike`)
 - 😊 “Shares statistical strength” between umpires
 - 😞 Heavily over-parametrized & indirect parametrization of `location`
 - 😞 Limited to two-way interactions (e.g. `umpire` \times `count`)
- Towards a fully nonparametric model
 - ▶ Uses Bayesian Additive Regression Trees
 - 😊 Avoid specifying functional form of F or potential interactions *a priori*
 - 🤖 Implementation challenge: handling all of the categorical inputs

Framing the problem

Estimating strike probabilities

Results & next steps

Model from D. & Wyner (2017)

- Let $y = 1$ for called strike and $y = 0$ for ball
- For umpire u , we model

$$\log \left(\frac{\mathbb{P}(y = 1)}{\mathbb{P}(y = 0)} \right) = \theta_0^u + \theta_b^u + \theta_c^u + \theta_p^u + \theta_{count}^u + \theta_{loc}^u \times \log \left(\frac{\hat{p}}{1 - \hat{p}} \right)$$

- Umpire-specific effects of batters, catchers, pitchers, & count
- For each catcher c : θ_c^u 's are i.i.d $\mathcal{N}(\bar{\theta}_c, \sigma_c^2)$
 - ▶ Shrink effects on each umpire to common mean $\bar{\theta}_c$
 - ▶ Common prior on $\bar{\theta}_c$'s shrinks all avg. catcher effects to 0
- \hat{p} is historical estimate of called strike prob. based on location

Model from D. & Wyner (2017)

- Let $y = 1$ for called strike and $y = 0$ for ball
- For umpire u , we model

$$\log \left(\frac{\mathbb{P}(y = 1)}{\mathbb{P}(y = 0)} \right) = \theta_0^u + \theta_b^u + \theta_c^u + \theta_p^u + \theta_{count}^u + \theta_{loc}^u \times \log \left(\frac{\hat{p}}{1 - \hat{p}} \right)$$

- Umpire-specific effects of batters, catchers, pitchers, & count
- For each catcher c : θ_c^u 's are i.i.d $\mathcal{N}(\bar{\theta}_c, \sigma_c^2)$
 - ▶ Shrink effects on each umpire to common mean $\bar{\theta}_c$
 - ▶ Common prior on $\bar{\theta}_c$'s shrinks all avg. catcher effects to 0
- \hat{p} is historical estimate of called strike prob. based on location
- ☹ 50+ hours to get 6,000 samples (fit in Stan in an HPC environment)
- ☹ Heavily over-parameterized ($\geq 170,000$ parameters)
- ☹ Only supported 2-way interaction (e.g. umpire \times catcher)

Bayesian Additive Regression Trees


- Chipman et al. (2010)'s probit model: $\mathbb{P}(y = 1) = \Phi(f(\mathbf{x}))$ where

$f(\mathbf{x}) = \text{sum of } M \text{ binary regression trees}$

- ☺ No need to specify functional form of f *a priori*
- ☺ BART has proven tremendously successful in applications
- Sums of regression trees have great representational flexibility

Gibbs Sampling

- Introduce latent $z \sim \mathcal{N}(f(\mathbf{x}), 1)$ s.t. $y = \mathbb{1}(z \geq 0)$
- Given sum of trees $f(\mathbf{x})$, draw z from truncated normal
- Given z , sequentially grow or prune each regression tree at random



Handling categorical predictors


- We have several categorical variables with many levels
 - ▶ Each season: 900+ batters, 100+ catchers, \sim 100 umpires

Handling categorical predictors

- We have several categorical variables with many levels
 - ▶ Each season: 900+ batters, 100+ catchers, ~ 100 umpires
- Default implementation of BART: create binary dummy variables

Default

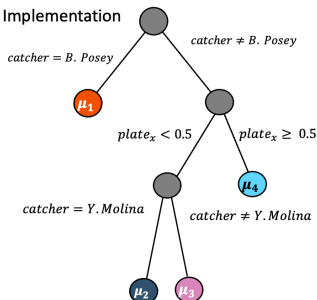
Implementation



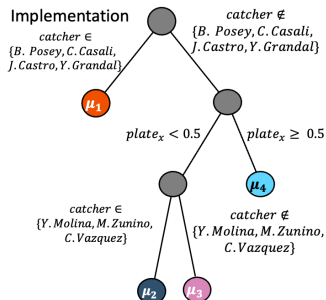
Handling categorical predictors

- We have several categorical variables with many levels
 - ▶ Each season: 900+ batters, 100+ catchers, ~ 100 umpires
- Default implementation of BART: create binary dummy variables
- Our implementation: randomly partition set of available levels

Default Implementation



Our New Implementation

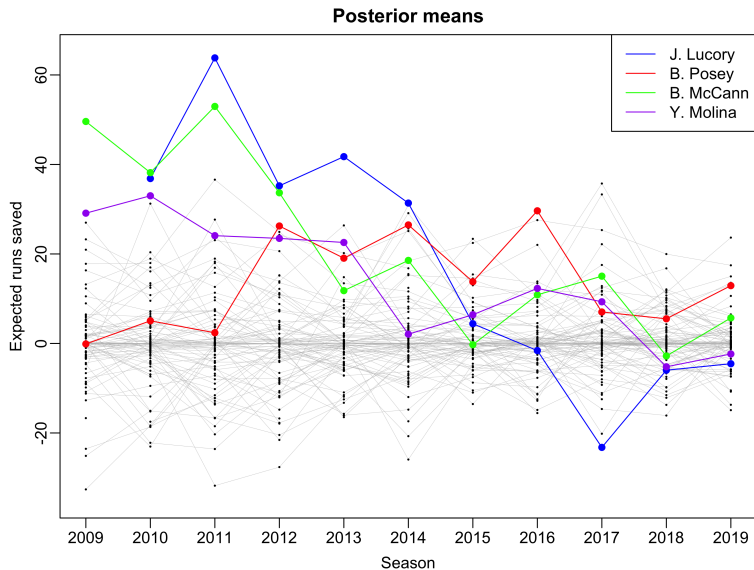


Framing the problem


Estimating strike probabilities

Results & next steps

Expected runs saved



Uncertainty in expected runs saved



New directions

- Causal estimates?
 - ▶ We summed $\hat{\rho} \times (\hat{\rho} - \hat{\rho}_0)$ for each catcher
 - Ⓜ Expected runs saved is confounded by pitcher, batter, location, ...
 - ▶ Partial salvage: average out location (see D. & Wyner (2017))

New directions

- Causal estimates?
 - ▶ We summed $\hat{\rho} \times (\hat{\rho} - \hat{\rho}_0)$ for each catcher
 - Ⓜ Expected runs saved is confounded by pitcher, batter, location, ...
 - ▶ Partial salvage: average out location (see D. & Wyner (2017))
- Smoother called strike estimates
 - ▶ For each (personnel, context) output a smooth heat map
 - ▶ Sum of treed GPs or sum of treed low rank thin plate splines?

New directions

- Causal estimates?
 - ▶ We summed $\hat{p} \times (\hat{p} - \hat{p}_0)$ for each catcher
 - Ⓜ Expected runs saved is confounded by pitcher, batter, location, ...
 - ▶ Partial salvage: average out location (see D. & Wyner (2017))
- Smoother called strike estimates
 - ▶ For each (personnel, context) output a smooth heat map
 - ▶ Sum of treed GPs or sum of treed low rank thin plate splines?
- Leveraging full pitch trajectory?
 - ▶ Effect of pitch movement & break on umpire's decision?
 - ▶ *Functional* inputs for treed models?

New directions

- Causal estimates?
 - ▶ We summed $\hat{\rho} \times (\hat{\rho} - \hat{\rho}_0)$ for each catcher
 - Ⓜ Expected runs saved is confounded by pitcher, batter, location, ...
 - ▶ Partial salvage: average out location (see D. & Wyner (2017))
- Smoother called strike estimates
 - ▶ For each (personnel, context) output a smooth heat map
 - ▶ Sum of treed GPs or sum of treed low rank thin plate splines?
- Leveraging full pitch trajectory?
 - ▶ Effect of pitch movement & break on umpire's decision?
 - ▶ *Functional* inputs for treed models?

Thanks, y'all!

Email: sameer.deshpande@wisc.edu

D. & Wyner (2017): [arXiv:1704.00823](https://arxiv.org/abs/1704.00823)

Website: <https://skdeshpande91.github.io>